ABSAAR ALI

Lead AI Engineer | Multi-Agent Systems & GenAI Architecture | LLM Orchestration

Location: Delhi, India
Contact: +91 90138 26608
Email: s.absaar.ali.job@gmail.com
LinkedIn: https://www.linkedin.com/in/syed-absaar-ali

Date of Birth: 3rd Oct 2001 **| Nationality:** Indian **| Gender:** Male

PROFESSIONAL SUMMARY

Lead AI Engineer with nearly two years of professional experience at Coforge, progressing through the ranks from Data Scientist → Senior AI Engineer → Lead AI Engineer. Specialising in multi-agent systems, GenAI architecture, and enterprise LLM orchestration across Banking, Aviation, and Financial Services domains. A published researcher recognised with the prestigious Miltos Petridis Memorial Trophy at the SGAI International Conference on Artificial Intelligence 2025 for pioneering work on aerodynamic surrogate modelling. Demonstrates deep expertise in agentic AI platforms (LangGraph, MCP, Azure OpenAI), full-stack delivery (Next.js, FastAPI, PostgreSQL), and finance-grade compliance (NIST AI RMF, EU AI Act, SR 11-7). Motivated by complex, enterprise-scale AI challenges and eager to contribute cutting-edge engineering expertise to organisations at the forefront of AI-driven transformation.

TECHNICAL SKILLS

- **Languages & Frameworks:** Python · TypeScript · C/C++ · JavaScript · SQL · PySpark · FastAPI · Next.js · React.js · Node.js
- **AI / ML & GenAI:** PyTorch · TensorFlow · Scikit-learn · XGBoost · CNNs · GANs · LangGraph · Multi-Agent Systems · LLM Orchestration · RAG · Prompt Engineering · Azure OpenAI (GPT-4o) · pyannote.audio · Socket.IO
- **Data & Cloud:** PostgreSQL · MySQL · MongoDB · FAISS · Neo4j · Azure (Document Intelligence, AI Search, Face API, Speech, OpenAI, Databricks, AKS, APIM) · Docker · OpenTelemetry · Git · Linux · PyTorch Ignite

PROFESSIONAL EXPERIENCE

Coforge, India
Data Scientist → Senior AI Engineer → Lead AI Engineer | July 2024 - Present

Lead AI Engineer | Multi-Agent Systems, GenAI Architecture, Technical Leadership Jan 2026– Present

Agentic AI Platform PoC — SDLC Automation for Northern Trust (MDP 2.0 / FOE)

- Directed three of eight workstreams (Input Validation, Context Assembly, Governance/Security/Observability) for an enterprise agentic AI PoC, targeting a 25-agent SDLC fleet (7 at MVP) on a shared LangGraph + Azure OpenAI + 13-server MCP infrastructure with 7-year WORM audit retention.
- Established and finalized cross-workstream contracts (Validated Request, Context Bundle, Policy Decision, AgentTrace/AuditEvent) and OpenTelemetry GenAI telemetry standards, preventing contract drift before production deployment.
- Engineered the typed Context Assembly layer: orchestrated retrievers over Azure AI Search + Neo4j with ACL pre-filtering, managed token-budget eviction in the 128k context window, and ensured cache-stable ordering.
- Conceived the Policy Decision Point abstraction and append-only audit projection; authored regulatory-mapping artifacts aligning with NIST AI RMF, EU AI Act, and SR 11-7, ensuring finance-grade compliance (99.95% availability, P95 <5s, RPO <5 min).

HiringLens AI — End-to-End L1 Interview Automation Platform (PoC to Production)

- Spearheaded full-stack implementation of HiringLens AI, automating first-round technical screening on Next.js 16, FastAPI, PostgreSQL 15, and Azure OpenAI GPT-4o with ≥90% selection relevance.
- Architected a three-channel Socket.IO system (control / answer / media) for interview lifecycle management, integrated live Azure Speech-to-Text transcription, and implemented concurrent biometric proctoring using Azure Face API and pyannote.audio speaker diarisation.
- Constructed an LLM evaluation pipeline (GPT-4o) with resume/JD-aware prompt templates, question-specific rubrics, and Semaphore-based rate limiting; produced structured STRONG_HIRE / HIRE / CONSIDER / NO_HIRE recommendations derived from AI scoring, communication metrics, and proctoring flags.

Senior AI Engineer | ML Engineering, Data Platform Modernisation Oct 2025 – Jan 2026

AI-Driven ETL Migration — Informatica & Greenplum to Databricks

- Led automated migration of legacy ETL pipelines from Informatica and Greenplum to Databricks, leveraging AI for pipeline translation, schema inference, and dependency sequencing.

- Implemented a graph-based dependency engine to ensure safe, lossless migration while preserving business logic integrity.
- Developed schema inference and AI-driven SQL dialect translation modules to auto-convert legacy transformations into PySpark-native Databricks notebooks, reducing manual rewrite effort by 60% and accelerating delivery.

Data Scientist | LLM Pipelines, Time-Series Forecasting, GenAI **July 2024 – Oct 2025**

ML-Driven Banking Price Optimisation — Savings Rate Forecasting

- Designed an XGBoost-based time-series model on 8 years of banking inflow data, leveraging rolling 6-month windows to predict 2–3 month demand horizons, enabling dynamic interest rate adjustments via Bayesian optimisation.

Dual-LLM Airway Bill Extraction — Azure Document Intelligence & GPT-4

- Developed a dual-LLM extraction workflow combining Azure Document Intelligence for OCR/layout analysis and GPT-4 for semantic field extraction and normalization.
- Implemented a human-in-the-loop feedback loop (HITL/RLHF) for low-confidence extractions, continuously refining model accuracy through operator corrections.

Samsung Research Institute Bangalore, India June 2023 – August 2023

Software Research & Development Intern | Node.js, Security, Android

- Designed and implemented a FIDO2-compliant passkey authentication system featuring a secure Relying Party server built with Node.js and WebAuthn, enabling seamless passwordless user sign-ins.
- Constructed REST APIs for WebAuthn credential registration and assertion flows, ensuring adherence to W3C security standards and compatibility across modern browser environments.
- Optimised metadata and database architecture for Android 14, achieving a 20% reduction in on-device storage consumption and enhancing data security through time-bound access controls.

RESEARCH & PUBLICATIONS

Car Drag Coefficient Prediction from 3D Point Clouds

SGAI International Conference on Artificial Intelligence (AI-2025) | Winner – Miltos Petridis Memorial Trophy | July 2025

- Invented a novel geometric surrogate model featuring a custom lightweight PointNet2D encoder paired with a Bi-LSTM temporal encoder, attaining state-of-the-art aerodynamic prediction accuracy ($MAE = 0.0061$, $R^2 = 0.9525$).
- Pioneered a slice-based architecture that processes 3D vehicle geometry as sequential 2D cross-sections, bypassing computationally expensive 3D convolutions and reducing aerodynamic analysis time from hours (CFD simulation) to seconds in early design phases.
- Built a complete MLOps pipeline — from dataset preprocessing and model training to live inference deployment — using PyTorch Ignite and Docker, ensuring reproducibility and operational scalability.

PROJECTS

AI Resume Standardisation Engine March 2025

- Built an LLM-powered system that converts unstructured resumes into structured, ATS-ready schemas with normalised skills taxonomies, chronological experience timelines, and standardised role classifications.
- Orchestrated multi-step agentic tool workflows encompassing document parsing, skill ontology mapping, vector retrieval, and ATS scoring to generate quantitative resume compatibility reports.

Low-Light Image Enhancement using Generative Adversarial Networks April 2024

- Implemented a GAN-based encoder–decoder architecture for paired low-light image enhancement, trained end-to-end on the LOLv2 benchmark dataset to recover brightness, contrast, and colour fidelity under challenging lighting conditions.
- Achieved approximately 17.2 dB PSNR and 0.78 SSIM on LOLv2 through systematic tuning of adversarial and reconstruction loss weightings and targeted data augmentation strategies for perceptually stable training.

ACHIEVEMENTS & CERTIFICATIONS

- **Miltos Petridis Memorial Trophy (2025):** Awarded for the best student paper in the Application Stream at the SGAI International Conference on Artificial Intelligence, organised by the British Computer Society.
- **GATE DA 2025 (Data Science and Artificial Intelligence):** Secured All India Rank 764 among more than 50,000 candidates in the Graduate Aptitude Test in Engineering.

EDUCATION

Delhi Technological University (Formerly Delhi College of Engineering), Delhi April 2024

Bachelor of Technology (B.Tech) — Major: Computer Science and Engineering | Minor: Machine Learning